

STAGE 3 PRACTICAL ASSESSMENT

Artist Intelligence & Recommendation Challenge

Build a small evidence-led system that understands demonstrated artist capabilities and recommends relevant artists from incomplete hirer briefs.

EXPECTED EFFORT Maximum 6 hours	SUBMISSION WINDOW 4 calendar days
DATASET 15 profiles · 4 briefs	

THE SITUATION

1. Two connected intelligence problems

A fictional creative marketplace has basic artist profiles, portfolio media and short hirer conversations. Hirer requests are usually incomplete. Artist profiles may also contain unsupported claims, missing media or evidence that is relevant to one kind of work but not another.

Your goal First, build an evidence-backed understanding of what each artist appears capable of. Then use that intelligence to produce useful initial recommendations from sparse hirer context and improve the ranking when new information arrives.

What you will receive

1. Four short hirer-conversation files. These are the only hirer-side inputs; there is no client-information packet.
2. 15 synthetic or anonymized artist folders: 5 photographers, 5 musicians and 5 video editors.
3. Each folder contains basic profile text and a small amount of category-relevant media: images, audio or short video. One case is incomplete or damaged.
4. One follow-up update for a hirer request, supplied separately, to test re-ranking.

Link: <https://drive.google.com/drive/folders/1yOAc9DaJ86oLReaSKdnATR3dHiCQyhwn?usp=sharing>

Before you build

Create decision_note.md before implementation. In no more than one page, state the decision your system supports, your first-version scope and non-goals, the capability dimensions you will use for each category, and the main assumptions or risks in your approach.

Scope boundary This is a controlled assessment, not unpaid production work. The data is synthetic or anonymised and the output has no production use. Do not build a frontend, scrape the web, train a model or deploy an application.

PART A

2. Artist intelligence

Create an evidence-backed capability record for every artist. You choose the detailed schema and category-specific dimensions, but the result must distinguish what the profile claims from what the supplied work actually demonstrates.

5. **Extract demonstrated capabilities.** Describe observable skills, style or genre, individual/team format and relevant experience signals.
6. **Use category-specific judgement.** The useful dimensions for a photographer, musician and editor should not be identical.
7. **Cite evidence.** Reference the source file and image identifier or timestamp for claims derived from media.
8. **Represent uncertainty.** State what cannot be determined and explain how confidence is produced.
9. **Select media sensibly.** Do not process every second or every frame blindly. Explain what you selected, skipped and why.

Important distinction Do not infer reliability, punctuality, character, popularity or professionalism from portfolio media or appearance. Those are trust signals and are outside this assignment.

PART B

3. Contextual recommendations

10. **Interpret the hirer's intent.** Separate explicit constraints, reasonable assumptions, contradictions and important unknowns.
11. **Recommend despite incompleteness.** Return two ranked artists whenever plausible matches exist; do not wait for a perfect brief.

12. **Connect capability to context.** Explain which demonstrated artist signals matter for this requirement and which do not.
13. **Show results before questions.** After the initial shortlist, add an Improve your matches section with no more than two questions.
14. **Prioritise questions.** Explain how each possible answer could materially change the ranking.
15. **Re-rank.** Process the supplied follow-up update and explain what changed and why.

Required outputs

File	Minimum content
artist_intelligence.jsonl	One record per artist: capability assessment, category-specific dimensions, evidence, unknowns and confidence
recommendations.json	Initial top two per brief, reasons, trade-offs, assumptions, uncertainty, then up to two refinement questions and expected impact
updated_recommendation.json	Revised ranking after the supplied follow-up, including what changed and why

SUBMISSION

4. What to submit

16. **Source code.** A private repository or ZIP with pinned dependencies and one documented run command.
17. **decision_note.md — maximum one page.** Framing, scope, category dimensions, non-goals and assumptions.
18. **README.md — maximum two pages.** Approach, setup, media selection, implemented choices, evaluation and limitations.
19. **Generated outputs.** The three files above; do not manually repair final results.
20. **AI_USAGE.md.** Which AI tools helped, what they produced and what you personally verified or changed.
21. **Demo — maximum five minutes.** Show one capability claim with evidence, one recommendation, the re-ranking and one weakness.

Rules

22. Timebox the work to six hours and report actual time spent. Prioritisation is part of the assessment.
23. Any reasonable combination of rules, LLMs, vision/audio models or embeddings may be used. Explain why it was appropriate.

24. Paid API usage is optional and capped at ₹300. More models, processing or spending earns no automatic advantage.
25. Do not identify or reverse-search people, hard-code IDs, publish the data, expose credentials or present manual edits as system output.
26. AI coding assistants are allowed. You remain responsible for every submitted claim, decision and line of code.

How we will assess it

Area	Points	What good looks like
Hirer intent understanding	15	Useful interpretation of incomplete conversations without demanding a perfect brief
Artist capability intelligence	15	Relevant, evidence-backed category-specific assessment without unsupported inference
Contextual recommendations	20	Capabilities are connected to the brief; rankings, questions and changes are defensible
Technical execution	10	Reproducible, selective, robust and efficient processing
Evaluation and uncertainty	10	Meaningful checks, honest limitations and safe treatment of missing/conflicting evidence
Live working discussion	30	Clarification, evidence defence, adaptation, feedback response and ownership

If shortlisted, you will join a 30-minute working discussion involving media evidence, a new sparse requirement, changed context and feedback on one part of your submission. No additional presentation is required.

What will not impress us A polished interface, elaborate architecture or confident artistic language cannot compensate for unsupported capability claims, generic rankings, hidden manual work or an inability to explain the submission.